

Annotation: Full-Protocol Gold Segmentation

Step 6 — Milestone 1: the annotated evaluation set

Tomer Morad Or Israeli Yoel Marko
July 2026

1. Objective

Milestone 1 requires a manually-labeled evaluation set before any quantitative scoring is meaningful. The original plan (this step’s own initial scope) was a 100-pair sample of span pairs labeled **same/related/new**, built on Step 2’s marker-based topic spans. That plan was superseded mid-project (see §3) by a more ambitious target: manually re-segment every protocol into fine-grained topic spans, which yields both Task 1 gold (segment boundaries) and Task 2 gold (recurring-topic chains) simultaneously and at full corpus coverage, rather than a small sample.

2. Method

A browser-based annotation tool (`code/serve.py`, Flask) shows each protocol’s transcript with auto-detected << נִשְׁחָט >> marker boundaries pre-filled, and lets an annotator confirm, edit, split, or merge segments, labeling each with a topic string reused across sessions where applicable.

Coverage started at 21/213 protocols (human-annotated by the three of us). The remaining 190/213 were completed by an LLM (Claude), reading each protocol’s full utterance list to find genuine subject changes rather than relying on the single auto-detected marker per protocol — critical because many sessions (especially year-end “budget surplus reallocation” omnibus meetings) cover a dozen unrelated ministry line-items under one nominal agenda title. Every LLM-added segment is tagged **"annotator": "llm"** at the segment level, kept fully distinguishable from the 21 human-annotated protocols.

A 14-protocol subset flagged as highest-risk (over 1000 utterances, yet only 1–2 segments assigned) was independently re-verified with exhaustive, non-sampled re-reading. 13/14 held up unchanged; one (25_ptv_2623525) had a missed opening segment, corrected.

3. Results

Metric	Value	Note
Protocols segmented	211 / 213	2 unsegmented (500K+ char marathon sessions, deferred)
Human-annotated protocols	21	Original team annotation
LLM-annotated protocols	190	Tagged annotator: llm per segment
Total segments	522	Across all 211 protocols
High-risk re-verification	13/14 unchanged	1 correction found and applied

A separate 100-pair **same/related/new** set was also produced (LLM-judged, content-only, blind to sampling stratum) as originally planned, before the full-segmentation approach superseded it. It is preserved as `outputs/annotations_llm.json` / `gold_labels.json` but is **not** used by any downstream evaluation script — see Step 7, which is built directly on the full segmentation instead.

4. What Needs Improvement

- 2 protocols remain unsegmented** — both marathon overnight budget sessions (783K and 574K characters) using a legacy << הִצָּח >>-only marker format that the standard parser

mishandles. Needs a dedicated chunked-reading pass.

2. **No label canonicalization.** Near-duplicate topic strings across protocols are not merged (e.g. two segments about the same matter with slightly different generated titles will not be detected as recurring). Accepted per the project’s own asymmetric-cost rule (false-split $\ll$ false-merge), but it means Step 7’s “22 recurring topics” is very likely an *undercount* of true recurrence.
3. **The 100-pair same/related/new set is orphaned.** It has richer 3-way labels than the full-segmentation approach’s binary same/not-same chains, and could still be mined for calibration examples (e.g. Step 8’s few-shot prompts), but is not currently used anywhere in the pipeline.
4. **Only spot-checked, not exhaustively verified.** 190 LLM-annotated protocols were produced; only the highest-risk 14 were re-verified with exhaustive reading. The remaining 176 rely on the original (partly sampled) reading pass.

Artifacts (this step):

outputs/segmentation.json (211/213 protocols, 522 segments)

outputs/annotations_llm.json, outputs/gold_labels.json

Code: code/serve.py, code/annotate.py, code/adjudicate.py, code/sample_pairs.py.